SYSTEMS ENGINEERING PORTFOLIO · LINUX / FOSS

41 programs, and something outside each one that decides if it is right.

Every project here is built against an **external oracle** — gcc and clang, the Linux kernel, `dash`, GNU tar, readelf, SQLite, Python's own regex engine. Not one of them is the judge of its own correctness. That constraint is the whole method, and it is why **10 of the 18 defects recorded here were in the test harness rather than in the code**.

paperpaladin
github.com/paperpaladin

41	926	52,736	35,927	41/41
REPOSITORIES	COMMITS	LINES OF SOURCE	LINES OF TESTS	SUITES GREEN

Every figure in this document is generated from the repositories themselves: commit counts from `git`, line counts from the files, test counts from running each suite. Nothing is transcribed by hand.

The method

A test you wrote, checking code you wrote, against a number you worked out yourself, proves that you were consistent. It does not prove you were right. **Every project here is measured against an authority it does not control.**



How these 41 compose. Each arrow is a dependency recorded in a PLAN.md. The three process tools in the bottom lane act on all the others: **unrepro** found a real reproducibility defect in eight of the 41, **commitq** gates every commit in all of them, and **srcdiff** audits published artifacts against the repositories they claim to come from.

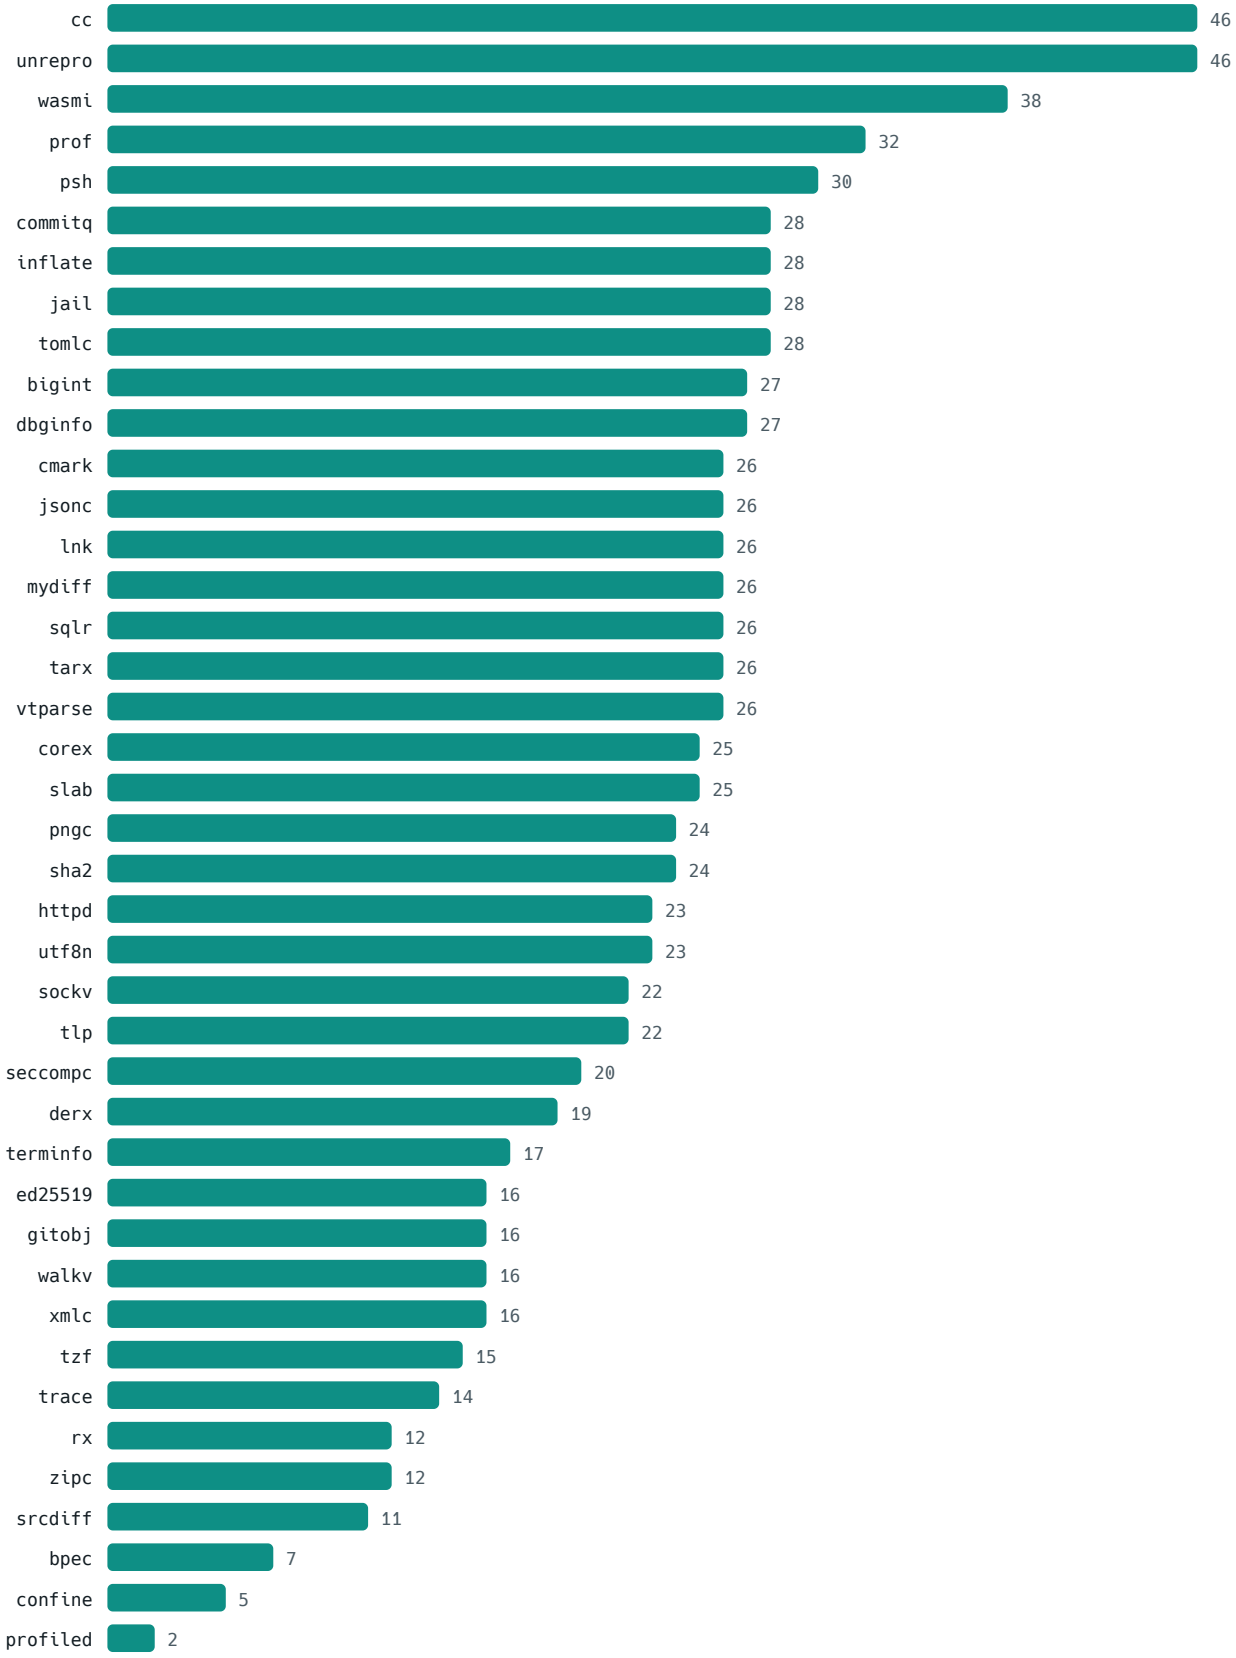
WHY THIS MATTERS MORE THAN IT SOUNDS

The first run of almost every harness in this project passed. The useful question is never “do my tests pass” but **“have I watched this harness fail?”** Where that question got asked, it usually turned out the answer was no — and the suite had been reporting success because it was silently doing nothing.



Defects recorded in commit messages and READMEs across all every repository here. Counted once per root cause.

Scale

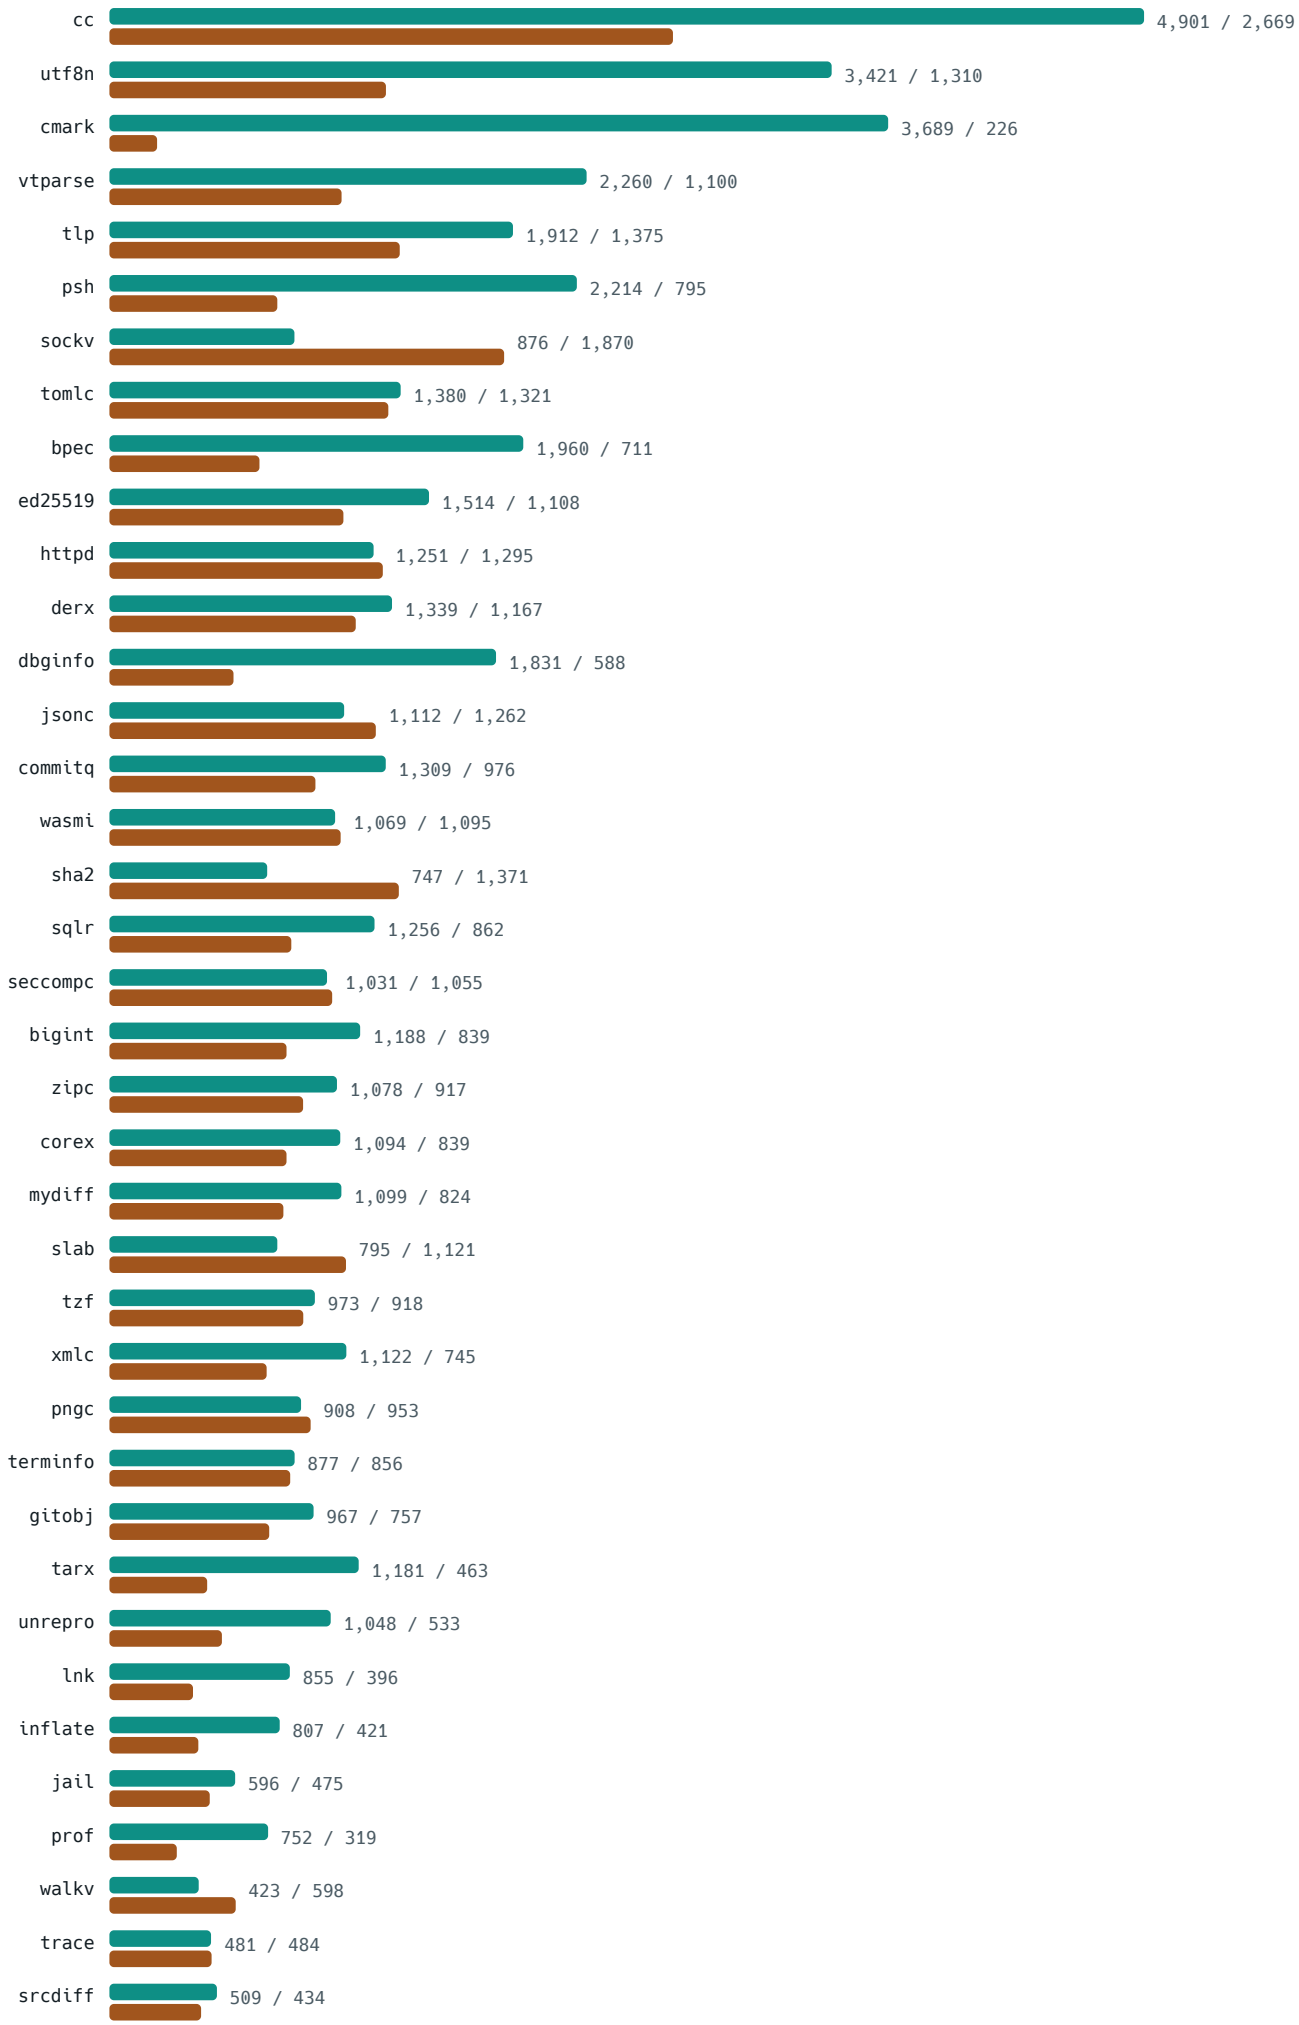


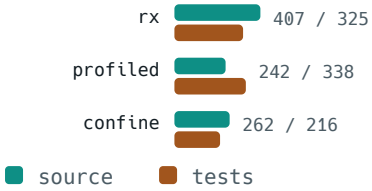
Commits per repository. Granularity is deliberate: every commit builds, verified by removing build artifacts and rebuilding each one in turn.

Source against tests

35,927 lines of tests to 52,736 of source. The ratio is the point: in several of these the harness **is** the project, and the program under test is

scaffolding for it.


cc
4,901 / 2,669
utf8n
3,421 / 1,310
cmark
3,689 / 226
vtparse
2,260 / 1,100
tlp
1,912 / 1,375
psh
2,214 / 795
sockv
876 / 1,870
tomlc
1,380 / 1,321
bpec
1,960 / 711
ed25519
1,514 / 1,108
httpd
1,251 / 1,295
derx
1,339 / 1,167
dbginfo
1,831 / 588
jsonc
1,112 / 1,262
commitq
1,309 / 976
wasmi
1,069 / 1,095
sha2
747 / 1,371
sqlr
1,256 / 862
seccompc
1,031 / 1,055
bigint
1,188 / 839
zipc
1,078 / 917
corex
1,094 / 839
mydiff
1,099 / 824
slab
795 / 1,121
tzf
973 / 918
xmlc
1,122 / 745
pngc
908 / 953
terminfo
877 / 856
gitobj
967 / 757
tarx
1,181 / 463
unrepro
1,048 / 533
lnk
855 / 396
inflate
807 / 421
jail
596 / 475
prof
752 / 319
walkv
423 / 598
trace
481 / 484
srcdiff
509 / 434



Lines of source and lines of test per repository, sorted by total. Labels read source / tests.

A NUMBER WORTH DISTRUSTING

Test *counts* are not comparable between these repos and are deliberately not charted together. `wasmi`'s 11,833 are generated arithmetic cases; `jail`'s 77 are hand-written assertions against kernel behaviour, and each one took longer to get right than a thousand of the former. Plotting them on one axis would imply a comparison that does not exist.

REPO	WHAT THE SUITE ACTUALLY REPORTS	PLAN
bigint	107573 agree, 0 disagree	20/28
bpec	36 passed, 0 failed	8/11
cc	434 passed, 0 failed (2 suites)	27/32
cmark	22 passed, 0 failed	18/18
commitq	156 passed, 0 failed	8/17
confine	20 passed, 0 failed	3/12
corex	72 passed, 0 failed, 0 skipped	17/22
dbginfo	85 passed, 0 failed	17/27
derx	216 passed, 0 failed, 0 skipped	11/22
ed25519	36 passed, 0 failed	9/12
gitobj	2146 passed, 0 failed	6/24
httpd	230 passed, 0 failed, 0 skipped	18/28
inflate	162 passed, 0 failed	21/30
jail	77 passed, 0 failed, 0 skipped	14/23
jsonc	366 passed, 0 failed	17/22
lnk	29 passed, 0 failed	18/31
mydiff	996 passed, 0 failed	13/17
pngc	74 passed, 0 failed	17/24
prof	15 passed, 0 failed, 0 skipped	18/30
profiled	17 passed, 0 failed	3/11
psh	212 agree, 0 disagree, 0 skipped	20/30
rx	3000 agree, 0 disagree, 0 skipped (seed=0)	5/16

REPO	WHAT THE SUITE ACTUALLY REPORTS	PLAN
seccompc	125 passed, 0 failed, 1 skipped	10/18
sha2	243 passed, 0 failed	18/23
slab	32 passed, 0 failed, 0 skipped	18/29
sockv	83 passed, 0 failed, 1 skipped	11/21
sqlr	123 agree, 0 disagree	17/28
srcdiff	30 passed, 0 failed	6/17
tarx	119 passed, 0 failed	7/19
terminfo	34 passed, 0 failed	7/8
tlp	130 passed, 0 failed	13/17
tomlc	630 passed, 0 failed	18/23
trace	43 passed, 0 failed	11/20
tzf	34 passed, 0 failed	5/13
unrepro	114 passed, 0 failed	12/24
utf8n	107 passed, 0 failed	7/17
vtparse	101 passed, 0 failed	9/20
walkv	349 scenarios checked, 0 failed	11/18
wasmi	11833 passed, 0 failed	25/35
xmlc	144 passed, 0 failed	8/18
zipc	76 passed, 0 failed	7/17

The repositories

Grouped by what kind of thing they have to survive. Each card names the authority that decides whether it is correct, and the one property

that, if violated, means the project failed.

Toolchain

Programs that build or run other programs.

cc C compiler for x86-64

46 commits · 4,901 lines · 2,669 test lines

cc → gcc AND clang must agree on every case

HARD INVARIANT	Self-host fixpoint: cmp cc2 cc3
RESULT	434 passed, 0 failed (2 suites)
PLAN	27 done, 5 queued

wasmi WebAssembly interpreter

38 commits · 1,069 lines · 1,095 test lines

wasmi → Expectations computed independently in Python

HARD INVARIANT	A malformed module is rejected or traps, never corrupts the host
RESULT	11833 passed, 0 failed
PLAN	25 done, 10 queued

psh POSIX shell subset

30 commits · 2,214 lines · 795 test lines

psh → dash: stdout AND exit status

HARD INVARIANT	Identical stdout and status to dash on every corpus script
RESULT	212 agree, 0 disagree, 0 skipped
PLAN	20 done, 10 queued

lnk Static ELF linker

26 commits · 855 lines · 396 test lines

lnk → The kernel runs the binary, or it does not

HARD INVARIANT	Every executable emitted runs, or we refused to emit it
RESULT	29 passed, 0 failed
PLAN	18 done, 13 queued

mydiff Myers diff and patch

26 commits · 1,099 lines · 824 test lines

mydiff → GNU diff byte-identical, and GNU patch applies our output

HARD INVARIANT	The loop closes both ways: our diff feeds GNU patch, GNU's diff feeds our -p
RESULT	`996 passed, 0 failed`
PLAN	13 done, 4 queued

gitobj git object database reader and writer

16 commits · 967 lines · 757 test lines

gitobj → The object name itself: git chose the filename, we recompute the SHA-1.

HARD INVARIANT	Every object read hashes back to its own name, loose or packed, through the whole delta chain
RESULT	`2146 passed, 0 failed`
PLAN	6 done, 18 queued

Formats & parsers

Programs that read a specification exactly.

inflate — DEFLATE / gzip codec

28 commits · 807 lines · 421 test lines

inflate → gzip -d accepts our output, both directions

HARD INVARIANT	gzip accepts all we emit; we decode all it emits
RESULT	162 passed, 0 failed
PLAN	21 done, 9 queued

tomlc — TOML 1.0.0 parser

28 commits · 1,380 lines · 1,321 test lines

tomlc → toml-test, plus Python's stdlib tomllib

HARD INVARIANT	Valid documents parse to the expected value, not merely without error
RESULT	630 passed, 0 failed
PLAN	18 done, 5 queued

dbginfo — DWARF debug-info parser

27 commits · 1,831 lines · 588 test lines

dbginfo → readelf --debug-dump

HARD INVARIANT	Every fact printed agrees with readelf
RESULT	85 passed, 0 failed
PLAN	17 done, 10 queued

cmark — CommonMark parser and HTML renderer

26 commits · 3,689 lines · 226 test lines

cmark → The specification's own spec.json: 652 exact markdown/HTML pairs

HARD INVARIANT	Output is byte-identical to the spec's expected HTML, nothing normalised
RESULT	22 passed, 0 failed
PLAN	18 done, 0 queued

jsonc — Strict RFC 8259 JSON parser

26 commits · 1,112 lines · 1,262 test lines

jsonc → JSONTestSuite: 95 must-accept, 188 must-reject, 35 pinned choices

HARD INVARIANT	A crash can never score as a rejection; exit status outside {0,1} is a failure
RESULT	366 passed, 0 failed
PLAN	17 done, 5 queued

sqlr SQLite file-format reader

26 commits · 1,256 lines · 862 test lines

sqlr → libsqlite3 itself

HARD INVARIANT — Byte-identical output to SQLite; malformed input errors, never crashes

RESULT — `123 agree, 0 disagree`

PLAN — 17 done, 11 queued

tarx tar reader and writer

26 commits · 1,181 lines · 463 test lines

tarx → GNU tar, bidirectional interop

HARD INVARIANT — GNU tar agrees both ways, and no member escapes the destination

RESULT — `119 passed, 0 failed`

PLAN — 7 done, 12 queued

vtparse VT500 parser and wcwidth

26 commits · 2,260 lines · 1,100 test lines

vtparse → The published DEC state table, all 3584 pairs

HARD INVARIANT — Never allocates, never reads past its byte, matches the table

RESULT — `101 passed, 0 failed`

PLAN — 9 done, 11 queued

pngc PNG decoder, all colour types and Adam7

24 commits · 908 lines · 953 test lines

pngc → PngSuite: 162 images that must decode, 14 that must be refused

HARD INVARIANT — A corrupt file is refused by CRC, never rendered as a wrong image

RESULT — `74 passed, 0 failed`

PLAN — 17 done, 7 queued

sha2 SHA-2 family, HMAC and PBKDF2

24 commits · 747 lines · 1,371 test lines

sha2 → NIST CAVP vectors, plus hashlib, coreutils and openssl agreeing

HARD INVARIANT — Byte-exact against published vectors; not constant-time, and says so first

RESULT — `243 passed, 0 failed`

PLAN — 18 done, 5 queued

utf8n Unicode normalisation, all four forms

23 commits · 3,421 lines · 1,310 test lines

utf8n → NormalizationTest.txt from the UCD: 806,840 assertions

HARD INVARIANT Normalising twice equals normalising once, for every form on every input

RESULT 107 passed, 0 failed

PLAN 7 done, 10 queued

tlp metamorphic SQL logic-bug finder

22 commits · 1,912 lines · 1,375 test lines

tlp → A relation SQLite must satisfy with itself, so no second engine is needed;

HARD INVARIANT Every disagreement reported is a real one, and a run that finds nothing is distinguishable from a run that tested nothing

RESULT 130 passed, 0 failed

PLAN 13 done, 4 queued

derx DER reader and X.509 printer

19 commits · 1,339 lines · 1,167 test lines

derx → openssl asn1parse and x509, as a diff

HARD INVARIANT Every element matches asn1parse byte for byte, or is refused with a message

RESULT 216 passed, 0 failed, 0 skipped

PLAN 11 done, 11 queued

terminfo compiled-terminfo reader

17 commits · 877 lines · 856 test lines

terminfo → ncurses as a subprocess: infocmp for every capability of all 44 terminals

HARD INVARIANT Every capability reported is the one infocmp reports, because a positional format can be parsed perfectly and named wrongly

RESULT 34 passed, 0 failed

PLAN 7 done, 1 queued

ed25519 Ed25519 and X25519, built up from the field

16 commits · 1,514 lines · 1,108 test lines

ed25519 → A second Ed25519 in Python written from the RFC algorithm rather thar

HARD INVARIANT Every layer is differenced against an implementation sharing no code with it, because sign-then-verify is not evidence

RESULT 36 passed, 0 failed

PLAN 9 done, 3 queued

xmlc — XML 1.0 well-formedness checker and canonicaliser

16 commits · 1,122 lines · 745 test lines

xmlc → The W3C XML Conformance Test Suite: 2585 catalogued cases, vendored

HARD INVARIANT: Acceptance is exactly the XML 1.0 Fifth Edition grammar, and no input crashes, overreads or leaks

RESULT: 144 passed, 0 failed

PLAN: 8 done, 10 queued

tzf — TZif timezone file reader

15 commits · 973 lines · 918 test lines

tzf → zdump and Python's zoneinfo must both agree

HARD INVARIANT: Every instant answered agrees with both readers; malformed input is refused

RESULT: 34 passed, 0 failed

PLAN: 5 done, 8 queued

zipc — ZIP archive reader and writer

12 commits · 1,078 lines · 917 test lines

zipc → Info-ZIP unzip 6.00 and zip 3.0, bidirectionally, across every method and

HARD INVARIANT: The central directory is authoritative, and no member can ever be written outside the destination

RESULT: 76 passed, 0 failed

PLAN: 7 done, 10 queued

bpec — Byte-level BPE, exact to the reference

7 commits · 1,960 lines · 711 test lines

bpec → The HuggingFace tokenizers library, token for token

HARD INVARIANT: The same token ids as the reference, in the same order

RESULT: 36 passed, 0 failed

PLAN: 8 done, 3 queued

Systems

Programs that must survive the machine misbehaving.

jail — Landlock sandbox

28 commits · 596 lines · 475 test lines

jail → The kernel: denied means EACCES, verified

HARD INVARIANT — An ungranted path is never reachable; the child cannot relax itself

RESULT — `77 passed, 0 failed, 0 skipped`

PLAN — 14 done, 9 queued

bigint — Arbitrary-precision integers

27 commits · 1,188 lines · 839 test lines

bigint → Python's native int, which is itself arbitrary precision

HARD INVARIANT — Every operation matches Python exactly, including floor division of negatives

RESULT — `107573 agree, 0 disagree`

PLAN — 20 done, 8 queued

slab — LD_PRELOAD malloc replacement

25 commits · 795 lines · 1,121 test lines

slab → Real programs must produce identical output

HARD INVARIANT — Two live allocations never overlap; every pointer aligned

RESULT — `32 passed, 0 failed, 0 skipped`

PLAN — 18 done, 11 queued

httpd — HTTP/1.1 static file server

23 commits · 1,251 lines · 1,295 test lines

httpd → curl as reference client, RFC 9112 for the error cases

HARD INVARIANT — No request can read a byte outside the document root

RESULT — `230 passed, 0 failed, 0 skipped`

PLAN — 18 done, 10 queued

sockv — netlink socket enumerator

22 commits · 876 lines · 1,870 test lines

sockv → ss from iproute2, both directions

HARD INVARIANT — Every socket the kernel reports is listed once with every field as given

RESULT — `83 passed, 0 failed, 1 skipped`

PLAN — 11 done, 10 queued

seccompc — seccomp-BPF policy compiler

20 commits · 1,031 lines · 1,055 test lines

seccompc → The kernel, plus an i386 probe driving the 32-bit bypass

HARD INVARIANT — A syscall the policy did not permit never runs, by any entry point

RESULT — `125 passed, 0 failed, 1 skipped`

PLAN — 10 done, 8 queued

walkv — Write-ahead-log key-value store

16 commits · 423 lines · 598 test lines

walkv → A crash injected at every I/O boundary

HARD INVARIANT — Recovered state is a prefix with no holes

RESULT — `349 scenarios checked, 0 failed`

PLAN — 11 done, 7 queued

rx — Linear-time regex engine

12 commits · 407 lines · 325 test lines

rx → Python's re module

HARD INVARIANT — A pc enters a thread list at most once per input position

RESULT — `3000 agree, 0 disagree, 0 skipped (seed=0)`

PLAN — 5 done, 11 queued

Observability

Programs that watch other programs without disturbing them.

prof — Sampling CPU profiler

32 commits · 752 lines · 319 test lines

prof → A planted hot loop with known 10:1 ratios

HARD INVARIANT	Profiling does not change the profiled program
RESULT	15 passed, 0 failed, 0 skipped
PLAN	18 done, 12 queued

corex — ELF core dump reader

25 commits · 1,094 lines · 839 test lines

corex → gdb: registers, backtrace and mappings compared value by value

HARD INVARIANT	Every value printed agrees with gdb, or it is not printed
RESULT	72 passed, 0 failed, 0 skipped
PLAN	17 done, 5 queued

trace — ptrace syscall tracer

14 commits · 481 lines · 484 test lines

trace → strace -c and strace -f -c

HARD INVARIANT	Tracing does not change the traced program
RESULT	43 passed, 0 failed
PLAN	11 done, 9 queued

confine — What a process can actually do

5 commits · 262 lines · 216 test lines

confine → A process built with unshare and setpriv, reported back exactly

HARD INVARIANT	A line says measured only if the operation was attempted this run
RESULT	20 passed, 0 failed
PLAN	3 done, 9 queued

profiled — Derive a seccomp policy by watching a program run

2 commits · 242 lines · 338 test lines

profiled → The kernel: a denied syscall kills the process, read from wait4

HARD INVARIANT	A syscall absent from the profile is killed, and the kernel says so
RESULT	17 passed, 0 failed
PLAN	3 done, 8 queued

Supply chain & process

Programs that check the other programs.

unrepro — Build-reproducibility differ

46 commits · 1,048 lines · 533 test lines

unrepro → Planted causes, each asserting its diagnosis

HARD INVARIANT	Never report clean on a build that is not reproducible
RESULT	114 passed, 0 failed
PLAN	12 done, 12 queued

commitq — Multi-repo commit gate and GUI

28 commits · 1,309 lines · 976 test lines

commitq → Throwaway repos in a temp directory

HARD INVARIANT	No commit is ever created with a date other than now
RESULT	156 passed, 0 failed
PLAN	8 done, 9 queued

srcdiff — Supply-chain artifact auditor

11 commits · 509 lines · 434 test lines

srcdiff → Real xz upstream artifacts

HARD INVARIANT	No false negatives in HIGH, few enough FPs to stay visible
RESULT	30 passed, 0 failed
PLAN	6 done, 11 queued

What went wrong in the harness

These are the defects worth reading. In each case the code looked fine, the suite was green, and **the thing that was broken was the thing meant to find breakage.**

`inflate` **`make test-asan` could not fail**

Reject tests assert exit code exactly 1; AddressSanitizer also exits 1. Found by deleting the classic `dist > o->n` bounds check and watching 146 tests pass while ASan printed a heap overflow to stderr. Fixing the exit code revealed that every malformed input had also been leaking 192KB, invisible for the same reason.

`walkv` **Half the crash matrix never crashed**

In torn-write mode the `fsync` boundary called the injection hook and discarded its return value, so every fsync scenario ran to completion and reported a perfect result.

`walkv` **Four planted defects, only one caught**

Mutation testing showed the matrix crashed, recovered, verified and then *stopped* — nothing ever wrote after a recovery, so a subtly wrong recovery was unobservable. Adding that one scenario took the suite from 83 to 163 and the defect now fails 20 of them.

`tarx` **Every symlink assertion was vacuous**

`diff -r` dereferences symbolic links, so the comparisons silently compared targets instead of links, and a dangling link aborted the diff outright.

`srcdiff` **Sixteen false positives on a clean release**

Against the genuine xz-5.4.5 tarball and its git tag, every HIGH finding was wrong — and each was an `m4/*.m4` file absent from git, the exact shape of the attack it exists to find. The real signal would have been one line in seventeen.

`dbginfo` **ASan's exit code collided with “malformed input”**

Both are 1, so the robustness sweeps would have scored a genuine heap overflow as a clean diagnostic.

`unrepro` **Reported REPRODUCIBLE because `make` did nothing**

Pointed at four sibling repos it compared zero outputs and called them clean. Zero outputs is now *inconclusive*.

`commitq` **The commit gate was not gating 12 of 18 repos**

The test command was a hardcoded table of six names; everything else returned success, so two thirds of the project could be committed without running anything.

`slab` **gcc deleted the oracle**

At `-O2` gcc knows what `calloc` promises, so it proved the zero-checking loop could not observe a non-zero byte and removed it. An allocator that never zeroed anything passed cleanly.

`cc` **Five hand-computed expectations were wrong**

Each cost a detour through “is this a codegen bug?” The compiler was right every time. `t/addcase.sh` now takes the expected value from gcc and refuses the case if clang disagrees.

What went wrong in the code

Found by the oracles, never by re-reading the source. Each of these produced plausible output and passed every check that existed at the time.

`rx` **The linear-time guarantee was quietly broken**

Injecting a start thread at every input position collided with the generation marks — and that dedup *is* the guarantee. The one bug that destroyed the project's only interesting property.

`lnk` **A binary the kernel refused to run**

File offset and virtual address were aligned independently. `BASE 0x400000` is divisible by every power of two up to 4MB, so they stayed in step by arithmetic luck until an 8MB-aligned array broke the congruence. The fix made the whole class unemittable rather than merely fixed.

`jail` **A perfect sandbox applied to the wrong process**

The rule-installing loop reused the argument parser's index, so `jail` built a flawless ruleset and then exec'd `--ro`. Every internal invariant held; only a real program reporting a real errno could see it.

`dbginfo` **Strings resolved to different real strings**

clang emits `DW_AT_name` as `strx1` before declaring `str_offsets_base` on the same DIE. Every lookup succeeded, no error, no marker — the second unit was simply named `main.c` where readelf said `lib.c`.

`trace` **A SIGTERM death reported as exit 0**

After passing a signal through, the tracee was already running, so the next `PTRACE_SYSCALL` hit `ESRCH`. Reintroduced once by a later refactor and caught again immediately by the same test.

`prof` **A race against the dynamic loader**

A one-shot `/proc/PID/maps` snapshot raced `ld.so`, truncating libc frames non-deterministically.

`sqlr` **A closing paren ended the column list**

Depth was decremented and tested in the same iteration, so the `)` of `VARCHAR(10,2)` looked like the end of the CREATE statement. `SELECT *` returned one column of four, with no error.

`srcdiff` **A funding link parsed as a repository**

`github.com/sponsors/hynek` resolved to owner *sponsors*. No offline test could reach it: the bug was in metadata resolution, and the tests all covered the comparison.

HOW THIS IS KEPT HONEST

Every commit builds. Verified by checking out each one in turn, deleting build artifacts so a stale binary cannot fake a pass, and rebuilding.

All 41 build reproducibly. Proven by `unrepro`, one of them, which found that eight of its siblings were not — `-g` writes the compiler's working directory into DWARF `.debug_line_str`. Every repo compiling with `-g` was affected and every repo without it was already clean, which is independent confirmation of the cause rather than a plausible guess.

No commit can be gated by a suite that does not run. `commitq` runs each repo's tests against the working tree before committing and refuses when they fail — or when a repo has no discoverable suite at all, which was a real bug it had for two thirds of the project.

143 places take a deliberately simple route, each marked `shortcut:` in the source and naming both the ceiling it accepts and the upgrade path. 528 planned steps are done and 344 remain, ordered, each recording why it is next.

Generated from the repositories by `.portfolio/build.py`. Data-fill colours validated against the paper surface with the dataviz palette checker: worst CVD ΔE 12.6, normal-vision ΔE 21.0, all six checks pass.